

PLONK with lookups

Ariel Gabizon
Aztec

Zachary J. Williamson
Aztec

March 3, 2020

1 Introduction

In many cases, SNARK circuit size is dominated by repeated operations, e.g. bitwise XOR or AND. We investigate an alternative approach, where for commonly used operations we precompute a look-up table of the legitimate (input, output) combinations.

In polynomial language, this ultimately boils down to proving polynomials are the same “up to multiplicities”. That is, suppose that the values in the look-up table are $\{t_i\}_{i \in [d]}$ and the values in the witness are $\{f_i\}_{i \in [n]}$. We show that the polynomials

$$F(X) := \prod_{i \in [n]} (X - f_i), G(X) := \prod_{i \in [d]} (X - t_i)$$

have the same roots, ignoring multiplicities, i.e. that for non-negative integers $\{e_i\}_{i \in [d]}$ we have $F(X) = \prod_{i \in [d]} (X - t_i)^{e_i}$. Bootle et. al [BCG⁺] gave an algorithm for this exact problem, also in the context of efficient SNARK arithmetization of a common operation (in their case, repeatedly checking that field elements correspond to certain convenient sparse representations of boolean strings).

Their algorithm requires committing to a vector of length $d \log n$ that contains for each $i \in [d], j \in [\log n]$ the value $t_i^{2^j}$. They also commit to the binary decomposition of the $\{e_i\}$, and using the two, prove that F is of the desired form.

We present here an arguably simpler protocol for the same problem, that doesn’t require explicitly representing the multiplicities. Let us use the notation $f \subset t$ as shorthand for $\{f_i\}_{i \in [n]} \subset \{t_i\}_{i \in [d]}$.

The idea is to look at a sorted version $\{s_i\}$ of the values $\{f_i\}$, and compare the set of non-zero *differences* in $\{s_i\}$ and $\{t_i\}$. Note that if $f \subset t$ we indeed have that these sets of differences are the same. However, the converse is not true: We can create a sequence of values $\{s_i\}$ having the same difference set as $\{t_i\}$, but with the differences appearing in *different order*; and in this case we won’t have $s \subset t$. As an illustrating example, take

$$t = \{1, 4, 8\}, s = \{1, 1, 4, 8, 8, 8\}, s' = \{1, 5, 5, 5, 8, 8\}$$

All three sets have the same difference set $\{3, 4\}$; but since those differences appear in different order in s' and t , we don't have $s' \subset t$. We also didn't address checking the starting point is the same in s and t . Both these issues can be solved by comparing *randomized* difference sets: We choose random $\beta \in \mathbb{F}$, and compare the non-zero elements in the sequences $\{t_i + \beta t_{i+1}\}_{i \in [d-1]}$, $\{s_i + \beta s_{i+1}\}_{i \in [n-1]}$.

We show that this check can be done efficiently using a “grand product argument” similar to the one used in [GWC19]’s permutation argument. A technicality, is that it is more convenient to assume *all* of t ’s values appear at least once in f . Which is why we in fact take s to be the sorted version of the concatenation of f and t .

Precise details follow in the next section.

2 The main scheme

Notation Fix integers n, d and vectors $f \in \mathbb{F}^n, t \in \mathbb{F}^d$. We will (ab)use the notation $f \subset t$ to mean $\{f_i\}_{i \in [n]} \subset \{t_i\}_{i \in [d]}$. Let $H = \{\mathbf{g}, \dots, \mathbf{g}^n\}$ be a subset of a multiplicative subgroup of order $n+1$. Denote by κ be an element of a non-trivial coset of H in \mathbb{F}^* . For a polynomial $f \in \mathbb{F}[X]$ and $i \in [n+1]$ we sometimes denote $f_i := f(\mathbf{g}^i)$. For a vector $f \in \mathbb{F}^n$, we also denote by f the polynomial in $\mathbb{F}_{<n}[X]$ with $f(\mathbf{g}^i) = f_i$.

When $f \subset t$, we say that f is *sorted by* t when values appear in the same order in f as they do in t . Formally, for any $i < i' \in [n]$ such that $f_i \neq f_{i'}$, if $j, j' \in [d]$ are such that $t_j = f_i, t_{j'} = f_{i'}$ then $j < j'$.

Now, given $t \in \mathbb{F}^d, f \in \mathbb{F}^n, s \in \mathbb{F}^{n+d}$, define bi-variate polynomials F, G as

$$F(\beta, \gamma) := (1 + \beta)^n \cdot \prod_{i \in [n]} (\gamma + f_i) \prod_{i \in [d-1]} (\gamma(1 + \beta) + t_i + \beta t_{i+1})$$

$$G(\beta, \gamma) := \prod_{i \in [n+d-1]} (\gamma(1 + \beta) + s_i + \beta s_{i+1})$$

we have

Claim 2.1. $F \equiv G$ if and only if

1. $f \subset t$, and
2. s is (f, t) sorted by t .

Proof. We write F, G as elements of $\mathbb{F}(\beta)[\gamma]$ while taking out a $(1 + \beta)$ factor as follows.

$$F(\beta, \gamma) = (1 + \beta)^{n+d-1} \cdot \prod_{i \in [n]} (\gamma + f_i) \prod_{i \in [d-1]} (\gamma + (t_i + \beta t_{i+1})/(1 + \beta))$$

$$G(\beta, \gamma) = (1 + \beta)^{n+d-1} \prod_{i \in [n+d-1]} (\gamma + (s_i + \beta s_{i+1})/(1 + \beta))$$

Suppose first that $f \subset t$ and $s \in \mathbb{F}^{n+d}$ is (f, t) sorted by t .

Then for each $j \in [d-1]$, there is an index $i \in [n+d-1]$ such that $(t_j, t_{j+1}) = (s_i, s_{i+1})$. The corresponding factors in F, G are equal. That is,

$$(\gamma + (t_j + \beta t_{j+1})/(1 + \beta)) \equiv (\gamma + (s_i + \beta s_{i+1})/(1 + \beta))$$

Let $P' \subset [n+d-1]$ the set of these $d-1$ indices i , and let $P := [n+d-1] \setminus P'$. The n indices $i \in P$ are such that $s_i = s_{i+1}$, and $\{s_i\}_{i \in P}$ equals $\{f_i\}_{i \in [n]}$. For each $i \in P$, the corresponding factor of G , will be

$$\gamma + (s_i + \beta s_{i+1})/(1 + \beta) = \gamma + s_i,$$

which, for the j such that $s_i = f_j$ will equal the factor $(\gamma + f_j)$ in F .

For the other direction, assume $F \equiv G$ as polynomials in $\mathbb{F}[\beta, \gamma]$. Then $F \equiv G$ also as elements of $\mathbb{F}(\beta)[\gamma]$. Since $\mathbb{F}(\beta)[\gamma]$ is a unique factorization domain, we know that the linear factors of F, G , as written above must be equal. Thus, for each $i \in [d-1]$, G must have a factor equal to $(\gamma + (t_i + \beta t_{i+1})/(1 + \beta))$. In other words, for some $j \in [n+d-1]$,

$$\gamma + (t_i + \beta t_{i+1})/(1 + \beta) = \gamma + (s_j + \beta s_{j+1})/(1 + \beta),$$

which implies $t_i + \beta t_{i+1} = s_j + \beta s_{j+1}$, and therefore $t_i = s_j, t_{i+1} = s_{j+1}$. Call $P' \subset [n+d-1]$ the set of these $d-1$ indices j . For any index $j \in [n+d-1] \setminus P'$, there must be a factor “coming from f ” in F that equals the corresponding factor in G . More precisely, for such j there exists $i \in [n]$ such that

$$\gamma + f_i = \gamma + (s_j + \beta s_{j+1})/(1 + \beta),$$

or equivalently

$$f_i + \beta f_i = s_j + \beta s_{j+1}$$

which implies $f_i = s_j = s_{j+1}$.

Thus, we know that whenever consecutive values in s are different, they are exactly equal to two consecutive values in t , and all values of f are values of t . □

Claim 2.1 motivates the following protocol. We use the ranged polynomial protocol terminology from [GWC19] to describe it.

It will be convenient to assume $d = n + 1$. (If $d \leq n$ pad t with $n - d + 1$ repetitions of the last element.)

Preprocessed polynomials: The polynomial $t \in \mathbb{F}_{<n}[X]$ describing the lookup values.

Inputs: $f \in \mathbb{F}_{<n}[X]$

Protocol:

1. Let $s \in \mathbb{F}^{2n+1}$ be the vector that is (f, t) sorted by t . Also denote by s the following polynomial in $\mathbb{F}_{<2n+2}[X]$ that represents s on $H \cup \kappa H$: $s(\mathbf{g}^i) = s_i$ for $i \in [n+1]$; and $s(\kappa \mathbf{g}^i) = s_{n+i}$ for each $i \in [n+1]$.
2. P_{poly} computes the polynomial s and sends it to \mathcal{I} .
3. V_{poly} chooses random $\beta, \gamma \in \mathbb{F}$ and sends them to P_{poly} .
4. P_{poly} computes a polynomial $Z \in \mathbb{F}_{<n+d}[X]$ that aggregates the value $F(\beta, \gamma)/G(\beta, \gamma)$ where F, G are as described above. Specifically, we let

- (a) $Z(\mathbf{g}) = 1$,
- (b) For $2 \leq i \leq n$

$$Z(\mathbf{g}^i) = \frac{(1 + \beta)^n \prod_{j < i} (\gamma + f_j) \cdot \prod_{1 \leq j < i} (\gamma(1 + \beta) + t_j + \beta t_{j+1})}{\prod_{1 \leq j < i} (\gamma(1 + \beta) + s_j + \beta s_{j+1}) (\gamma(1 + \beta) + s_{n+j} + \beta s_{j+1})},$$

and

- (c) $Z(\mathbf{g}^{n+1}) = 1$.

5. P_{poly} sends Z to \mathcal{I} .
6. V_{poly} checks that Z is indeed of the form described above, and that $Z(\mathbf{g}^{n+d}) = 1$. More precisely, V_{poly} checks the following identities on H .

- (a) $L_1(Z - 1) \equiv 0$
- (b) $Z(X)(1 + \beta) \cdot (\gamma + f(X))(\gamma(1 + \beta) + t(X) + \beta t(X \cdot \mathbf{g})) \equiv Z(X \cdot \mathbf{g})(\gamma(1 + \beta) + s(X) + \beta s(X \cdot \mathbf{g}))(\gamma(1 + \beta) + s(\kappa \cdot X) + \beta s(\kappa \cdot X \cdot \mathbf{g}))$.
- (c) $L_1(s(\kappa X \cdot \mathbf{g}) - s(X/\mathbf{g})) \equiv 0$
- (d) $L_n(Z(X \cdot \mathbf{g}) - 1) \equiv 0$

and outputs acc iff all checks hold.

Lemma 2.2. *Suppose that $\{f(\mathbf{g}^i)\}_{i \in [n]}$ is not contained in $\{t(\mathbf{g}^i)\}_{i \in [d]}$. Then for any strategy of \mathcal{A} playing the role of P_{poly} in the above protocol, the probability that V_{poly} accepts is $\text{negl}(\lambda)$.*

Proof. Using Claim 2.1 we know that when f 's range is not contained in t 's for any choice of s sent by P_{poly} the polynomials $F(\beta, \gamma), G(\beta, \gamma)$ are different. From the SZ lemma e.w.p $\text{negl}(\lambda)$ V_{poly} chooses β, γ such that $F(\beta, \gamma) \neq G(\beta, \gamma)$. In this case we have $Z(\mathbf{g}^{n+d}) \neq 1$ which means V_{poly} rejects. \square

3 Generalizing to vector lookups and multiple tables

Suppose we have several witness polynomials $f_1, \dots, f_w \in \mathbb{F}_{<n}[X]$, and a table of values $t^* \in (\mathbb{F}^w)^d$. We wish to check that for each $j \in [n]$ $(f_1(\mathbf{g}^j), \dots, f_w(\mathbf{g}^j)) \in t^*$. We can use randomization to efficiently reduce to the case of Section 2.

For each $i \in [w]$ we will include in the set of preprocessed polynomials $t_i \in \mathbb{F}_{<d}[X]$ with $t_i(\mathbf{g}^j) = t_{i,j}$ for each $j \in [d]$.

The verifier will choose random $\alpha \in \mathbb{F}$.

Then we will define $t := \sum_{i \in [w]} \alpha^i t_i$, $f := \sum_{i \in [w]} \alpha^i f_i$.

Assume that for some $j \in [n]$, $(f_1(\mathbf{g}^j), \dots, f_w(\mathbf{g}^j)) \notin t^*$. Then e.w.p $d \cdot w/|\mathbb{F}|$, $f(\mathbf{g}^j) \notin t$. Thus, after the selection of α , we can run the protocol of the previous section on f, t .

3.1 multiple tables

Suppose further that we have in fact have multiple tables t_1^*, \dots, t_ℓ^* and for each $i \in [n]$ wish to check that for some predefined $j \in [\ell]$,

4 A simplified solution for continuous ranges

References

- [BCG⁺] J. Bootle, A. Cerulli, J. Groth, S. K. Jakobsen, and M. Maller. Arya: Nearly linear-time zero-knowledge proofs for correct program execution. In *Advances in Cryptology - ASIACRYPT 2018 - 24th International Conference on the Theory and Application of Cryptology and Information Security, Brisbane, QLD, Australia, December 2-6, 2018, Proceedings, Part I*, volume 11272 of *Lecture Notes in Computer Science*, pages 595–626.
- [GWC19] A. Gabizon, Z. J. Williamson, and O. Ciobotaru. PLONK: permutations over lagrange-bases for oecumenical noninteractive arguments of knowledge. *IACR Cryptology ePrint Archive*, 2019:953, 2019.